



# Setting up LLMs on CURC Resources

# Setting up LLMs on CURC Resources

## Instructor: Brandon Reyes

- Research Computing
- Website: [www.rc.colorado.edu](http://www.rc.colorado.edu)
- Documentation: <https://curc.readthedocs.io>
- Helpdesk: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Survey: <http://tinyurl.com/curc-survey18>



# Meet the User Support Team



Layla  
Freeborn



Brandon  
Reyes



Andy  
Monaghan



Michael  
Schneider



John  
Reiland



Dylan  
Gottlieb



Mohal  
Khandelwal



Ragan  
Lee

# Slides

[https://github.com/ResearchComputing/setting\\_up\\_llms\\_on\\_curc\\_resources\\_rc\\_short\\_course](https://github.com/ResearchComputing/setting_up_llms_on_curc_resources_rc_short_course)



# Session Overview

- What are LLMs?
- What is an LLM framework?
- Accessing CURC hosted LLMs
- Ollama
- Transformers by Hugging Face

# What are LLMs?

Large Language Models (LLMs) are types of artificial intelligence models that are trained on large amounts of data, which enables them to “understand” and generate natural language. They are great at:

- Summarizing material
- Constructing code (if they are trained for coding tasks)
- Conversational tasks
- Generating text

When you break it down, LLMs are made possible by multiplication and non-linear transformations. This allows them to run very quickly on GPUs.

# Important Considerations for LLMs

- Not all LLMs are equal, some are trained for specific tasks, such as code generation
- The type of data the LLMs are trained on determines what they can do
- General models (ChatGPT, Gemini) are great for common knowledge tasks, but can fall short in certain technical areas
- Bigger is not always better
  - You can have small LLMs that are extremely performant for a certain task
- LLMs can “hallucinate”. Put simply, this means they can generate incorrect answers or content that may sound plausible
  - Don’t treat LLMs as if they are all-knowing! They are not ... yet

# Proprietary vs Open-Source LLMs

- Models such as ChatGPT are proprietary and require that you run the model on their dedicated hardware
  - Pros: these companies have a lot of resources, allowing for quicker access and their models are usually state-of-the-art
  - Cons: **Data privacy and cost**
- Open-source LLMs can be ran anywhere
  - Pros: you can run them for free on our resources, your data stays on our system, and there are some very capable open-source models
  - Cons: proprietary models are not going to be available and complex LLMs can consume a lot of GPU memory

# What is an LLM framework?

Now that we know why we would like to run an LLM locally, how do we do that? In the past, running an LLM locally took quite a bit of effort. However, now there are several LLM frameworks available.

An LLM framework allows for a user-friendly way to access LLMs and utilize them

- Can enable access to proprietary or local LLMs
- Can enable running the LLM from the command line or within a script

# Common LLM Frameworks

There are many LLM frameworks out there, some with their own unique functionality. However, in this talk we will focus on two extremely popular LLM frameworks that enable you to run models locally:

- Ollama
- Transformers by Hugging Face

# What is Quantization and Instruct?

- Quantization
  - When training, most LLMs use either single or double precision. This can make those models consume a lot of memory.
  - Quantization allows you to reformat the model into a lower-precision format, which makes the model consume less memory.
- Instruct
  - Instruct or instruction models are ideal for specifying tasks for the LLM to perform and are the models used in common chat interfaces.

# Before we jump into things ...

Before running or working with LLMs on CURC resources, please be sure that you are adhering to all [CURC User Policies](#) and CU's policies and guidance around Artificial Intelligence outlined in the pages [Artificial Intelligence at CU Boulder](#) and [Resources & Guidance](#).

# Accessing CURC hosted LLMs

To streamline access and reduce redundant storage of LLMs, CURC has created a shared space for common LLMs

- This is a CURC managed space, you cannot store models in this space
  - If you do think a model should be here, reach out to us!
- All models can be accessed using the environment variable `CURC_LLM_DIR`
  - You must be on a compute node to get access to this variable and the provided LLMs

# CURC hosted Ollama Compatible LLMs

- gpt-oss:20b
  - An open-source model from OpenAI (ChatGPT folks)
  - Requires about 14 GB of GPU memory
- gemma3:12b
  - An open-source model from Google (a derivative of Gemini)
  - Requires around 8 GB of GPU memory
- llama3.1:8b
  - An open-source model from Meta
  - Requires about 5 GB of GPU memory

All of these models are quantized, instruct models, and available at:

`$CURC_LLM_DIR/ollama`

# CURC hosted Transformer Compatible LLMs

- gpt-oss-20b
  - This model has been quantized
  - Requires about 14 GB of GPU memory
- gemma-3-12b-it
  - This model has not been quantized
  - Without quantization, requires more than 20 GB of GPU memory
- Llama-3.1-8B-Instruct
  - This model has not been quantized
  - Without quantization, requires more than 20 GB of GPU memory

All of these models are instruct models and available at:

`$CURC_LLM_DIR/hf-transformers`

# Ollama

Ollama is an LLM framework maintained by Meta

- It is open-source and extremely beginner friendly
- Enables users to run LLMs locally and retrieve models
- Easily installable models are curated by Ollama
  - Available models can be seen at <https://ollama.com/search>
  - You can use other models, but you need to convert them to a compatible format
- With Ollama, you establish a server and then use this server to install, manage, and run LLMs

# Using CURC's Ollama install

The latest version of Ollama is available via our module stack. Once on a compatible compute node, you can load the module using:

`module load ollama`

- Starts an Ollama server for you and sets all necessary environment variables
- Uses CURC's hosted LLMs by default

If you would like to point to your own directory that contains LLMs, this can be done with:

```
export OLLAMA_MODELS=/projects/$USER/my_ollama_models; module load  
ollama
```

This module is only available on CPU nodes and NVIDIA GPU nodes. We only provide the most up-to-date version, and these versions will be updated during each planned maintenance

# Useful Ollama Commands

- **ollama list**
  - Lists all available models
- **ollama pull <model>**
  - Downloads a model if it does not exist, otherwise updates the model
- **ollama run <model>**
  - Runs the model from the command line and downloads the model if it does not exist
- **ollama rm <model>**
  - Removes the specified model

# Accessing Ollama in Python

To use Ollama within Python, you must install the appropriate packages. We have created a minimal uv environment for this. You can obtain access to this environment using:

```
module load uv
```

```
source $CURC_UV_ENV_DIR/ollama-python-api-env/bin/activate
```

[Note: The Ollama server needs to be running before you run your code.](#)

# Example Ollama Python script

[ollama\\_test.py](#)

```
from ollama import chat
from ollama import ChatResponse

response: ChatResponse = chat(model='llama3.1:8b', messages=[
    {
        'role': 'user',
        'content': 'In one sentence, how cool is CU Research Computing?',
    },
])
print(response.message.content)
```

To run this script from the command line:

(ollama-python-api-env) \$ [python ollama\\_test.py](#)

# Installing Ollama

As noted, our modules only install the newest version of Ollama that is available during our planned maintenance. Sometimes you need an older or newer version. We provide extremely detailed instructions for installing both Ollama and setting up the Ollama Python package via uv in the “Self-install instructions” tab:

<https://curc.readthedocs.io/en/latest/software/llms.html#ollama>

# Transformers by Hugging Face

Transformers is an LLM framework maintained by Hugging Face and is extremely popular

- It is open-source and contains additional functionality that Ollama does not have, such as the ability to train models
- Enables users to run LLMs locally and retrieve models
- Does not require you to setup a server like Ollama

# Transformers by Hugging Face

Models are not as easy to install as Ollama

- Pros:
  - Provides access to cutting-edge LLMs and datasets: <https://huggingface.co/>
  - Provides access to community driven models and models without quantization
- Cons:
  - Most well-known models require a Hugging Face account and the acceptance of terms of use to download
  - If the model uses a different library to run, you have to install those dependencies
  - Community driven models and datasets are sometimes not vetted

# Let's take a look at a model card!

<https://huggingface.co/google/gemma-3-12b-it>

# Using CURC's Transformers install

The latest version of Transformers is available via our module stack. Once on a compatible compute node, you can load the module using:

`module load hf-transformers`

- Loads base packages using uv and sets general environment variables
- Unlike Ollama, Transformers isn't connected to CURC's LLM directory (more on this later)

This module is only available on CPU nodes and NVIDIA GPU nodes. We only provide the most up-to-date version, and these versions will be updated during each planned maintenance

# Installing Transformers

As noted, our modules only install the newest version of Transformers that is available during our planned maintenance. Sometimes you need an older or newer version. We provide extremely detailed instructions for installing Transformers via uv in the “Self-install instructions” tab:

<https://curc.readthedocs.io/en/latest/software/llms.html#transformers-by-hugging-face>

# Downloading models

- You must have a Hugging Face account
- You need to generate a user access token

Once your account is setup, you can install models using:

```
hf download --local-dir <install-directory> <hf-model>
```

For in-depth instructions, see

<https://curc.readthedocs.io/en/latest/software/llms.html#downloading-models>

# Running an LLM using Transformers

# Streaming to the command line

# Adding Quantization

# Thank you!

Survey and feedback

<http://tinyurl.com/curc-survey18>



## Slides

[https://github.com/ResearchComputing/setting\\_up\\_llms\\_on\\_curc\\_resources\\_rc\\_short\\_course](https://github.com/ResearchComputing/setting_up_llms_on_curc_resources_rc_short_course)

